

Support Vector Machine

Classification

Linear classifiers: how would you separate these two groups of training examples with a straight line?


Feature
vector

Class value



$$f(x, w, b) = \text{sign}(w \circ x + b)$$

$$= \text{sign}(\sum w[i]x[i] + b)$$



Linear classifiers: how would you separate these two groups of training examples with a line?

Feature
vector



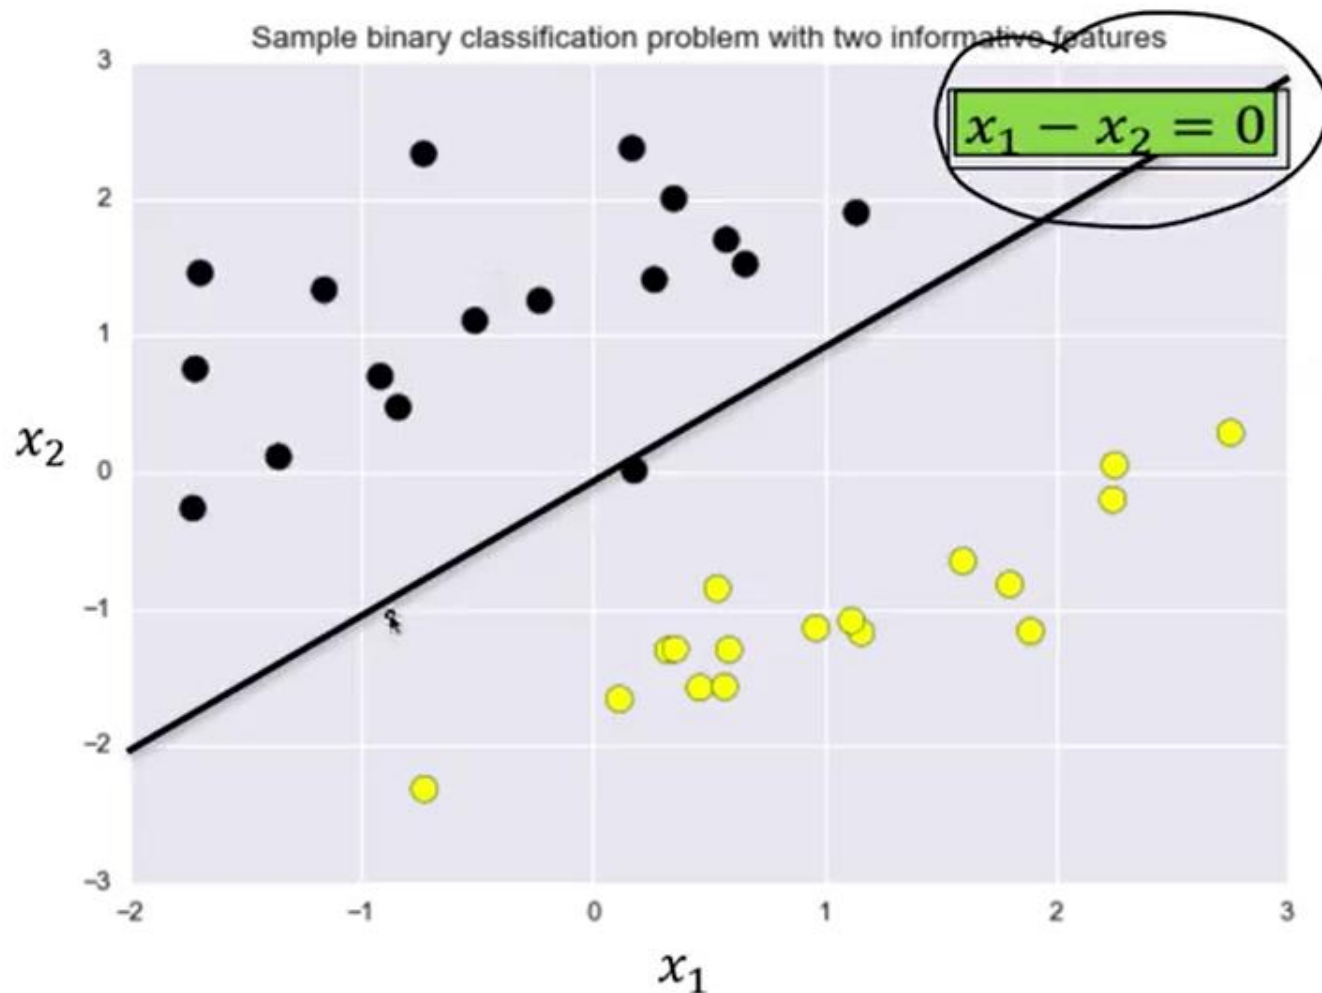
Class value

$$f(x, w, b) = \text{sign}(w \circ x + b)$$

$$x_1 - x_2 = 0$$

$$w = [1, -1]$$

$$b = 0$$



Linear classifiers: how would you separate these two groups of training examples with a line?

Feature
vector

Class value



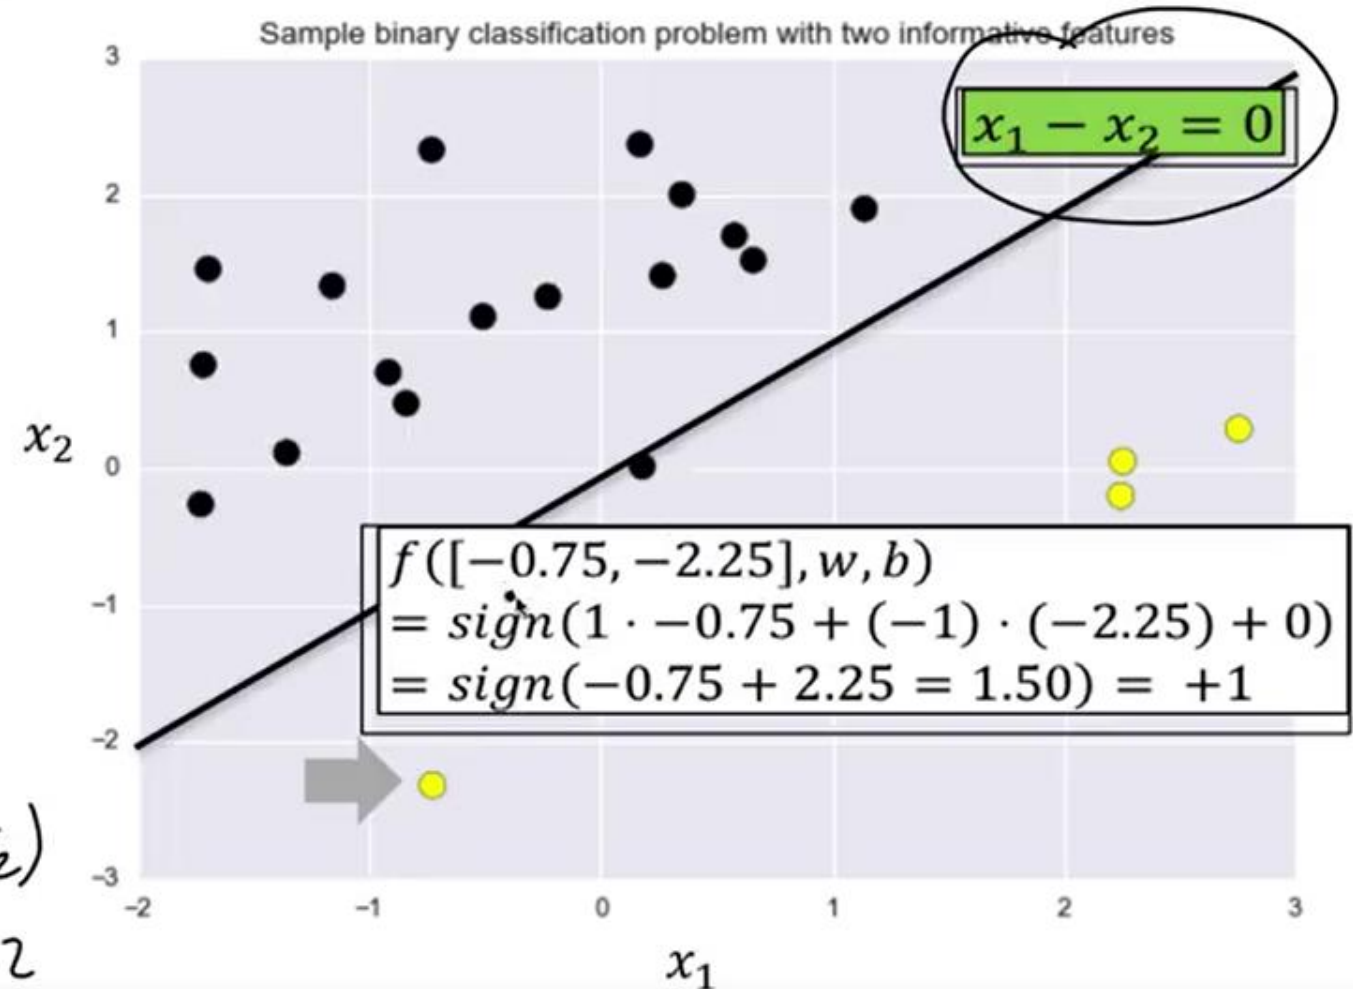
$$f(x, w, b) = \text{sign}(w \circ x + b)$$

$$x_1 - x_2 = 0$$

$$w = [1, -1]$$

$$b = 0 \quad (v_1, w_2) \circ (x_1, x_2) \\ = w_1 x_1 + w_2 x_2$$

$$= x_1 - x_2$$



Linear classifiers: how would you separate these two groups of training examples with a line?


Feature
vector

Class value



$$f(x, w, b) = \text{sign}(w \circ x + b)$$

$$\begin{aligned}x_1 - x_2 &= 0 \\ w &= [1, -1] \\ b &= 0\end{aligned}$$




Classifier Margin



$$f(x, w, b) = \text{sign}(w \circ x + b)$$

Classifier margin

Defined as the maximum width the decision boundary area can be increased before hitting a data point.




Classifier Margin



$$f(x, w, b) = \text{sign}(w \circ x + b)$$

Classifier margin

Defined as the maximum width the decision boundary area can be increased before hitting a data point.



Maximum Margin Linear Classifier: Linear Support Vector Machines




$$f(x, w, b) = \text{sign}(w \circ x + b)$$

Maximum margin classifier


The linear classifier with maximum margin is a linear Support Vector Machine (LSVM).

Or a SVM with linear kernel



Linear Support Vector Machine

[illegible]



Regularization for SVMs: the C parameter


- The strength of regularization is determined by C
- Larger values of C : less regularization
 - *Fit the training data as well as possible*
 - *Each individual data point is important to classify correctly*
- Smaller values of C : more regularization
 - *More tolerant of errors on individual data points*

```
In [21]: from sklearn.svm import LinearSVC
from adspy_shared_utilities import plot_class_regions_for_classifier

X_train, X_test, y_train, y_test = train_test_split(X_C2, y_C2,
                                                    random_state = 0)
fig, subaxes = plt.subplots(1, 2, figsize=(9, 3))

for this_C, subplot in zip([0.00001, 100], subaxes):
    clf = LinearSVC(C=this_C).fit(X_train, y_train)
    title = 'Linear SVC, C = {:.5f}'.format(this_C)
    plot_class_regions_for_classifier_subplot(clf, X_train, y_train,
                                             None, None, title, subplot)
```

Figure 7



Linear Models: Pros and Cons

Pros:


- Simple and easy to train.
- Fast prediction.
- Scales well to very large datasets.
- Works well with sparse data.
- Reasons for prediction are relatively easy to interpret.

Cons:


- For lower-dimensional data, other models may have superior generalization performance.
- For classification, data may not be linearly separable (more on this in SVMs with non-linear kernels)

SVM – Multi-class Classification


Binary classification:







Multi-class classification:





One-vs-all (one-vs-rest):




Class 1:  

Class 2:  

Class 3:  

$$\underline{h_{\theta}^{(i)}(x)} = P(y = i | x; \theta) \quad (i = 1, 2, 3)$$



Multi-class Classification with Linear Models


```
clf = LinearSVC(C=5, random_state = 67)
clf.fit(X_train, y_train)

print(clf.coef_)

[[-0.23401135  0.72246132]
 [-1.63231901  1.15222281]
 [ 0.0849835   0.31186707]
 [ 1.26189663 -1.68097    ]]

print(clf.intercept_)

[-3.31753728  1.19645936 -2.7468353  1.16107418]
```



Multi-class Classification with Linear Models

```
clf = LinearSVC(C=5, random_state = 67)
clf.fit(X_train, y_train)
```

```
print(clf.coef_)
```

```
[[-0.23401135  0.72246132]
 [-1.63231901  1.15222281]
 [ 0.0849835   0.31186707]
 [ 1.26189663 -1.68097    ]]
```


```
print(clf.intercept_)
```

```
[-3.31753728  1.19645936 -2.7468353  1.16107418]
```

```
y_apple = -0.23401135 * height + 0.72246132 * width - 3.31753728
```


```
height=2, width=6: y_apple = + 0.549 (>= 0: predict apple)
```

```
height=2, width=2: y_apple = - 2.340 (< 0: predict other)
```




Kernelized SVM

We saw how linear support vector classifiers could effectively find a decision boundary with maximum margin




Easy for a linear classifier

But what about more complex binary classification problems?




Easy for a linear classifier




Difficult/impossible for a linear classifier


A simple 1-dimensional classification problem for a linear classifier




A more perplexing 1-d classification problem for a linear classifier



Let's transform the data by adding a second dimension/feature
(set to the squared value of the first feature)




The data transformation makes it possible to solve this with a linear classifier




What does the linear decision boundary correspond to in the original input space?

Original input space

Feature space




Example of mapping a 2D classification problem to a 3D feature space to make it linearly separable



$$\mathbf{x}_i = (x_0, x_1)$$

$$\mathbf{v}_i = (x_0, x_1, 1 - (x_0^2 + x_1^2))$$

Transforming the data can make it much easier for a linear classifier.



Original input space

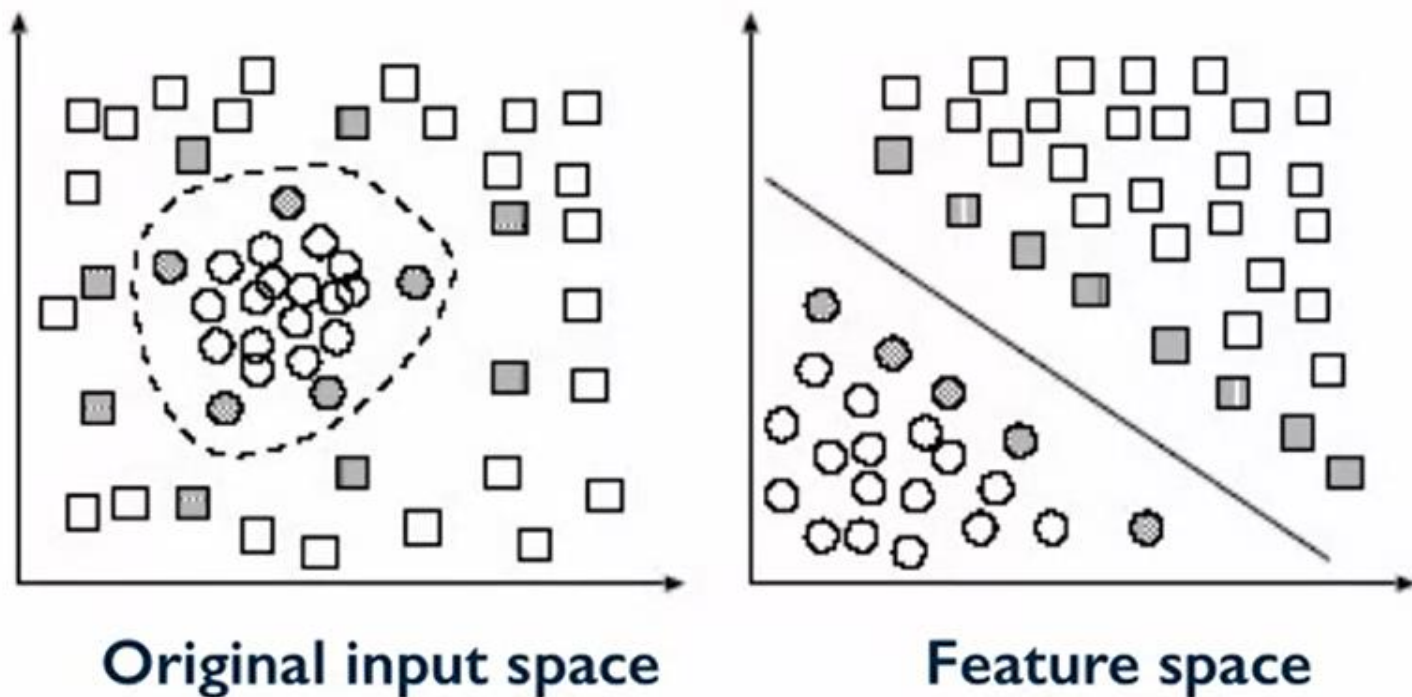
Feature space

Source: Wikipedia "Kernel Machine" article.

<https://commons.wikimedia.org/w/index.php?curid=47868867>


Radial Basis Function Kernel

$$K(\mathbf{x}, \mathbf{x}') = \exp [-\gamma \cdot \|\mathbf{x} - \mathbf{x}'\|^2]$$




A kernel is a similarity measure (modified dot product) between data points


Applying the SVM with RBF kernel




Support Vector Classifier: Linear Kernel



Support Vector Classifier: RBF kernel



Support Vector Classifier: Polynomial kernel, degree = 3



Radial Basis Function kernel: Gamma Parameter


$$K(\mathbf{x}, \mathbf{x}') = \exp [-\gamma \cdot \|\mathbf{x} - \mathbf{x}'\|^2]$$



gamma (γ): kernel width
parameter


small gamma (0.01)

large gamma (10)



Squared distance
between \mathbf{x} and \mathbf{x}'

The effect of the RBF gamma parameter on decision boundaries



Support Vector Machine with RBF kernel: using both C and gamma parameter

```
In [ ]: from sklearn.svm import SVC
        from adspy_shared_utilities import (
            plot_class_regions_for_classifier_subplot)

        from sklearn.model_selection import train_test_split

        X_train, X_test, y_train, y_test = train_test_split(X_D2, y_D2,
                                                            random_state = 0)
        fig, subaxes = plt.subplots(3, 4, figsize=(15, 12))

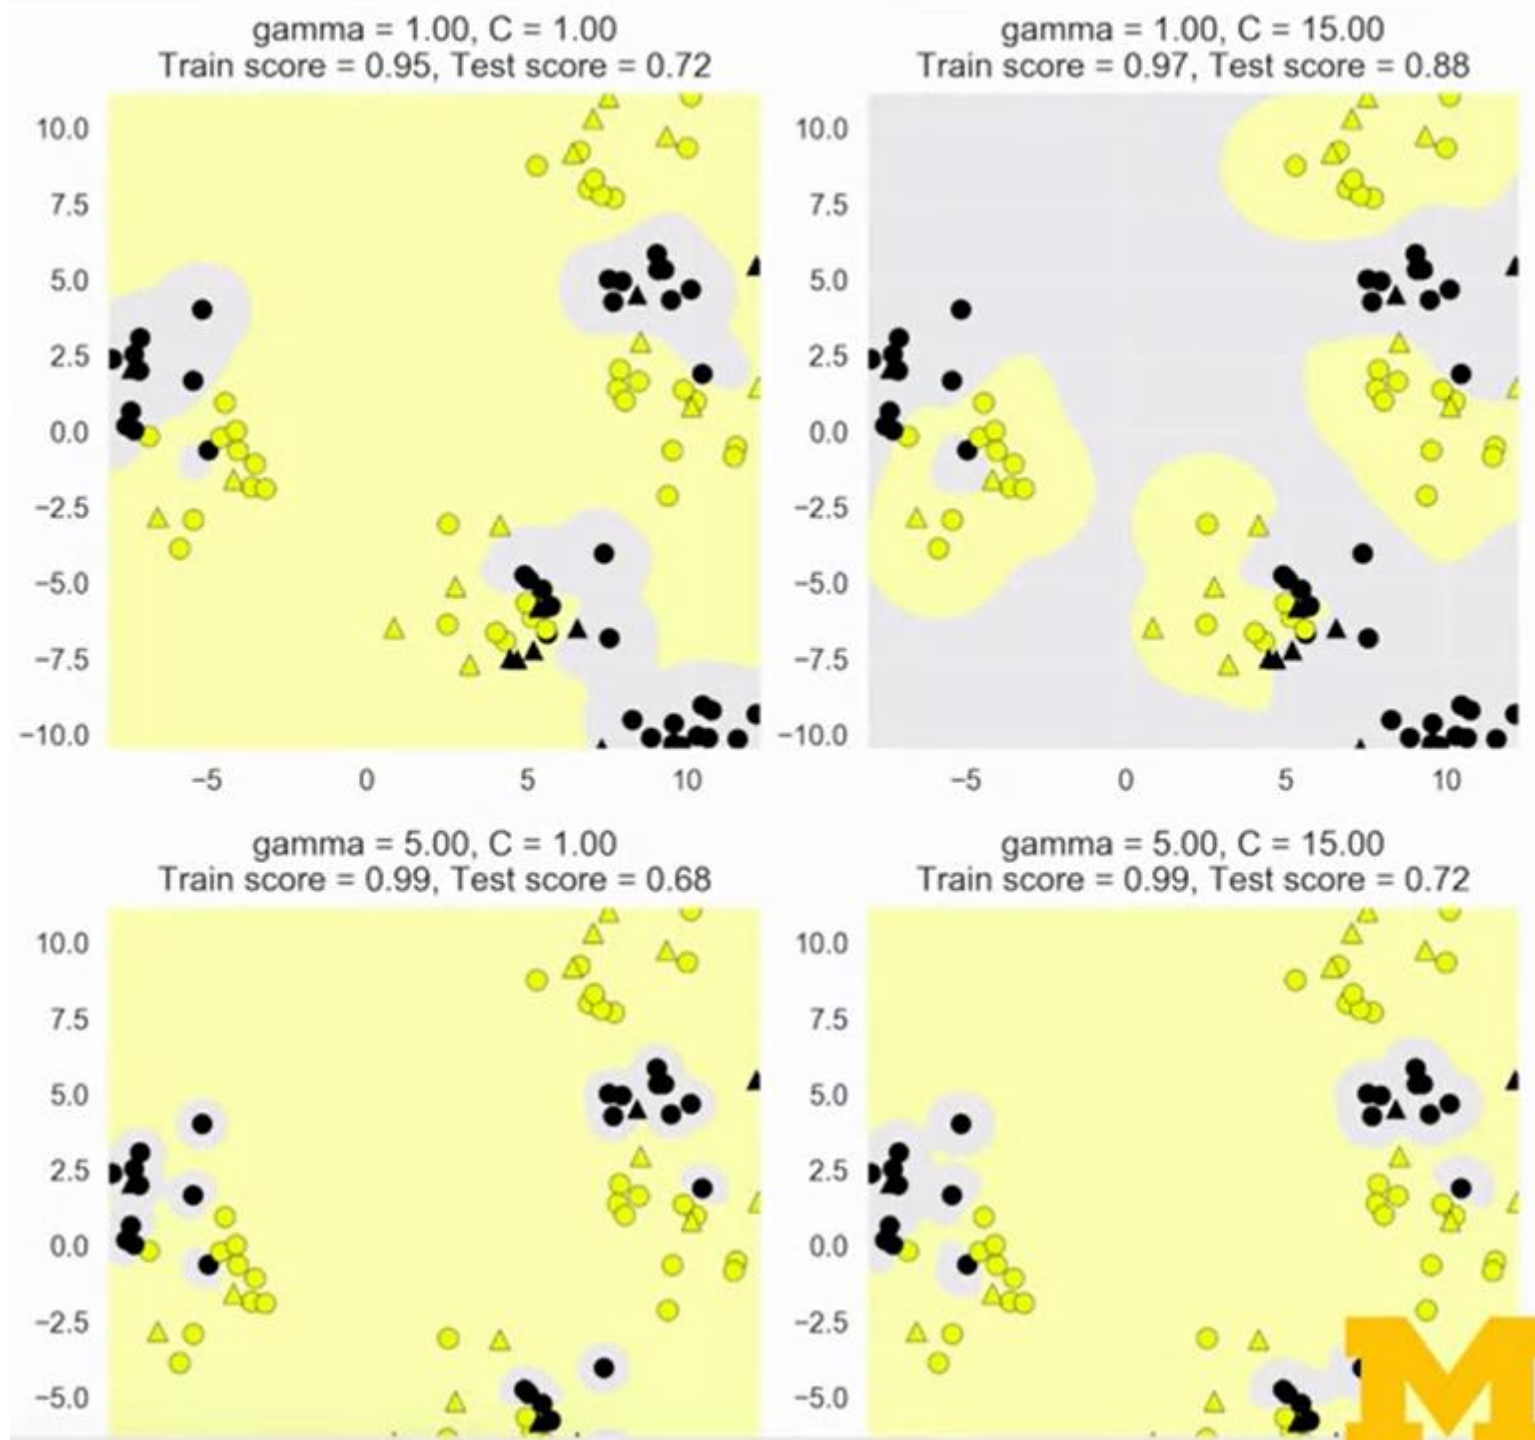
        for this_gamma, this_axis in zip([0.01, 1, 5], subaxes):

            for this_C, subplot in zip([0.1, 1, 15, 250], this_axis):
                title = 'gamma = {:.2f}, C = {:.2f}'.format(this_gamma, this_C)
                clf = SVC(kernel = 'rbf', gamma = this_gamma,
                          C = this_C).fit(X_train, y_train)
                plot_class_regions_for_classifier_subplot(clf, X_train, y_train,
                                                         X_test, y_test, title,
                                                         subplot)

            plt.tight_layout(pad=0.4, w_pad=0.5, h_pad=1.0)
```



- Using both
- Gamma and C parameter



Kernelized Support Vector Machines: pros and cons

Pros:

- Can perform well on a range of datasets.
- Versatile: different kernel functions can be specified, or custom kernels can be defined for specific data types.
- Works well for both low- and high-dimensional data.

Cons:

- Efficiency (runtime speed and memory usage) decreases as training set size increases (e.g. over 50000 samples).
- Needs careful normalization of input data and parameter tuning.
- Does not provide direct probability estimates (but can be estimated using e.g. Platt scaling).
- Difficult to interpret why a prediction was made.

Kernelized Support Vector Machines (SVC): Important parameters

Model complexity

- **kernel:** Type of kernel function to be used
 - *Default = 'rbf' for radial basis function*
 - *Other types include 'polynomial'*
- **kernel parameters**
 - *gamma (γ): RBF kernel width*
- **C:** regularization parameter
- Typically C and gamma are tuned at the same time.